

# BasicPremiereLeagueTablePrediction

2022-10-13

The first step in this Premier League table prediction is to read in historical data over the last 5 seasons.

```
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.2 --
## v ggplot2 3.3.6      v purrr  0.3.5
## v tibble  3.1.8      v dplyr  1.0.10
## v tidyr   1.2.1      v stringr 1.4.1
## v readr   2.1.3      v forcats 0.5.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()

library(ggplot2)

pl22 = read.csv('/Users/Shane/OneDrive/Desktop/Project-Work-main/Git/PremiereLeague/PL21-22.csv')
pl21 = read.csv('/Users/Shane/OneDrive/Desktop/Project-Work-main/Git/PremiereLeague/PL20-21.csv')
pl20 = read.csv('/Users/Shane/OneDrive/Desktop/Project-Work-main/Git/PremiereLeague/PL19-20.csv')
pl19 = read.csv('/Users/Shane/OneDrive/Desktop/Project-Work-main/Git/PremiereLeague/PL18-19.csv')
pl18 = read.csv('/Users/Shane/OneDrive/Desktop/Project-Work-main/Git/PremiereLeague/PL17-18.csv')

premlge = merge(merge(merge(merge(
  pl22,
  pl21, all = T),
  pl20, all = T),
  pl19, all = T),
  pl18, all = T)
premlge$Date = as.POSIXct(premlge$Date, format = "%m/%d/%y")

premlge = premlge %>%
  select(Date, HomeTeam, AwayTeam, FTHG, FTAG, FTR)
```

Next, I found the overall league scoring average over the last 5 seasons, as well as each teams individual home and away scoring averages.

```
lgAvg = data.frame(avgHG = mean(premlge$FTHG), avgAG = mean(premlge$FTAG))

tmAvgH = premlge %>%
  group_by(HomeTeam) %>%
  summarize(GoalsFor = mean(FTHG),
            GoalsAgainst = mean(FTAG))
tmAvgA = premlge %>%
  group_by(AwayTeam) %>%
```

```

summarize(GoalsFor = mean(FTAG),
          GoalsAgainst = mean(FTHG))
tmAvgH

```

```

## # A tibble: 28 x 3
##   HomeTeam      GoalsFor GoalsAgainst
##   <chr>          <dbl>         <dbl>
## 1 Arsenal        2.01           1.03
## 2 Aston Villa    1.40           1.51
## 3 Bournemouth    1.37           1.49
## 4 Brentford      1.16           1.11
## 5 Brighton       1.09           1.32
## 6 Burnley        1.01           1.31
## 7 Cardiff        1.11            2
## 8 Chelsea        1.76           0.884
## 9 Crystal Palace 1.16           1.25
## 10 Everton        1.4            1.23
## # ... with 18 more rows

```

From these averages, I create a power index to rank a teams goals for and goals against rating against the league average. For offense, anything above 1 is better than the league average (i.e. scores more), and for defense anything less than 1 is better (i.e. holds teams to less goals scored).

```

tmHPwr = tmAvgH %>%
  mutate(#Team = HomeTeam,
         OPwrH = GoalsFor/lgAvg$avgHG,
         DPwrH = GoalsAgainst/lgAvg$avgAG)

tmAPwr = tmAvgA %>%
  mutate(#Team = AwayTeam,
         OPwrA = GoalsFor/lgAvg$avgAG,
         DPwrA = GoalsAgainst/lgAvg$avgHG)

tmHPwr

```

```

## # A tibble: 28 x 5
##   HomeTeam      GoalsFor GoalsAgainst OPwrH DPwrH
##   <chr>          <dbl>         <dbl> <dbl> <dbl>
## 1 Arsenal        2.01           1.03  1.34  0.825
## 2 Aston Villa    1.40           1.51  0.938 1.21
## 3 Bournemouth    1.37           1.49  0.915 1.19
## 4 Brentford      1.16           1.11  0.774 0.884
## 5 Brighton       1.09           1.32  0.732 1.05
## 6 Burnley        1.01           1.31  0.675 1.04
## 7 Cardiff        1.11            2    0.739 1.60
## 8 Chelsea        1.76           0.884 1.17  0.707
## 9 Crystal Palace 1.16           1.25  0.774 1.00
## 10 Everton        1.4            1.23  0.936 0.985
## # ... with 18 more rows

```

Next, I created a list of this years current PL teams, substituting Norwhich for Nottingham since Nottingham hasn't been to the Premier League in over 20 years. Norwhich is a team relegated from last season, whose skill level is comparable to Norwhich.

```

teams = c(unique(tmHPwr$HomeTeam))
plTeams = c('Arsenal', 'Aston Villa', 'Bournemouth', 'Brentford',
            'Brighton', 'Chelsea', 'Crystal Palace', 'Everton',
            'Fulham', 'Leeds', 'Leicester', 'Liverpool',
            'Man City', 'Man United', 'Newcastle', 'Southampton',
            'Tottenham', 'West Ham', 'Wolves', 'Norwich')
# Substituting Norwich for Nottingham

tmHPwrPL = tmHPwr %>%
  filter(HomeTeam %in% plTeams) %>%
  select(HomeTeam, OPwrH, DPwrH)

tmAPwrPL = tmAPwr %>%
  filter(AwayTeam %in% plTeams) %>%
  select(AwayTeam, OPwrA, DPwrA)

```

Now I will load in the matches for this upcoming season, and bind each matchup to the respective teams home and away scoring power rankings.

```

matches = read.csv('/Users/Shane/OneDrive/Desktop/Project-Work-main/Git/PremiereLeague/PLMatches.csv')

dfT = data.frame(matrix(ncol = 6, nrow = 0))
colnames(dfT) = c('HomeTeam', 'AwayTeam', 'OPwrA', 'DPwrA', 'OPwrH', 'DPwrH')

for (i in plTeams) {
  m = matches %>%
    filter(HomeTeam == i)
  df1 = merge(m, tmHPwrPL, by = 'HomeTeam')
  dfT1 = merge(df1, tmAPwrPL, by = 'AwayTeam')

  m2 = matches %>%
    filter(AwayTeam == i)
  df2 = merge(m2, tmAPwrPL, by = 'AwayTeam')
  dfT2 = merge(df2, tmHPwrPL, by = 'HomeTeam')

  dfT = rbind(dfT, dfT2, dfT1)
}

dfT = dfT %>%
  distinct()

```

Now I will estimate the scores of each matchup based off of each teams power rankings.

- Estimated Home Goals = Home O Power Ranking \* Away D Power Ranking \* average goals scored in league
- Estimated Away Goals = Away O Power Ranking \* Home D Power Ranking \* average goals scored in league

```

scores = dfT %>%
  mutate(EstHGoals = OPwrH * DPwrA * lgAvg$avgHG,
         EstAGoals = DPwrH * OPwrA * lgAvg$avgAG)
scoreEst = scores %>%

```

```

select(HomeTeam, AwayTeam, EstHGoals, EstAGoals)

fixtures = scoreEst %>%
  select(home = HomeTeam, away = AwayTeam)

head(fixtures)

```

```

##           home    away
## 1   Aston Villa Arsenal
## 2   Bournemouth Arsenal
## 3    Brentford Arsenal
## 4    Brighton Arsenal
## 5     Chelsea Arsenal
## 6 Crystal Palace Arsenal

```

In the world of soccer, it is commonly accepted that the amount of goals scored comes from a poisson distribution. From this, I take the expected goals for each match from above, and use that value as the average in a poisson distribution. This allows me to find the probabilities of each team scoring anywhere from 0-7 goals against each other.

```

maxGoals = 7

allPredictions = map2_df(
  scoreEst$EstHGoals, scoreEst$EstAGoals,
  function(lambdaH, lambdaA, maxGoals) {
    hgoalProb = dpois(0:maxGoals, lambdaH) %>% `names<-`(0:maxGoals)
    agoalProb = dpois(0:maxGoals, lambdaA) %>% `names<-`(0:maxGoals)
    outer(hgoalProb, agoalProb) %>%
      as.data.frame() %>%
      gather() %>%
      rownames_to_column('row') %>%
      mutate(hgoal = as.numeric(row) %%(maxGoals + 1) - 1) %>%
      mutate(hgoal = case_when(hgoal < 0 ~ maxGoals, TRUE ~ hgoal),
             agoal = as.numeric(key)) %>%
      select(sample_hgoal = hgoal, sample_agoal = agoal, prob = value)
  },
  maxGoals) %>%
  cbind(fixtures[rep(seq_len(nrow(fixtures)), each = (maxGoals+1)^2), ], .) %>%
  group_by(home, away) %>%
  mutate( prob = prob/sum(prob) ) %>%
  ungroup()

head(allPredictions, 8)

```

```

## # A tibble: 8 x 5
##   home      away sample_hgoal sample_agoal      prob
##   <chr>    <chr>         <dbl>         <dbl>    <dbl>
## 1 Aston Villa Arsenal         0         0 0.0499
## 2 Aston Villa Arsenal         1         0 0.0685
## 3 Aston Villa Arsenal         2         0 0.0470
## 4 Aston Villa Arsenal         3         0 0.0215
## 5 Aston Villa Arsenal         4         0 0.00738
## 6 Aston Villa Arsenal         5         0 0.00202

```

```
## 7 Aston Villa Arsenal      6      0 0.000463
## 8 Aston Villa Arsenal      7      0 0.0000908
```

Since this data frame is very large, I created a nested list of each individual matchup and the probabilities of potential scores

```
nestedProb = allPredictions %>%
  nest(probabilities = c(sample_hgoal, sample_agoal, prob))

nestedProb2 = nestedProb %>%
  mutate(sampled_result = map(probabilities, sample_n, 1, weight = prob)) %>%
  select(-probabilities) %>%
  unnest(cols = c(sampled_result))

head(nestedProb2)
```

```
## # A tibble: 6 x 5
##   home      away  sample_hgoal sample_agoal  prob
##   <chr>    <chr>      <dbl>        <dbl>  <dbl>
## 1 Aston Villa Arsenal      2          1 0.0764
## 2 Bournemouth Arsenal      0          1 0.0846
## 3 Brentford  Arsenal      0          2 0.0695
## 4 Brighton   Arsenal      0          2 0.0835
## 5 Chelsea     Arsenal      4          0 0.0252
## 6 Crystal Palace Arsenal      2          1 0.0723
```

nestedProb2 is to display the score that is most likely to happen between two teams.

Now, I use a monte carlo method to simulate the entirety of the Premier League season n times. This allows the model to simulate matches over and over to try to predict the overall most likely match results over an entire season. While doing this, I assigned the match results to the Premier League Format of 3 points for a win, 1 for a tie, and 0 for a loss.

```
n = 1000
matchSims = rerun(n, nestedProb %>%
  mutate(sampled_result = map(probabilities, sample_n, 1, weight = prob)) %>%
  select(-probabilities) %>%
  unnest(cols = c(sampled_result)) %>%
  select(-prob) %>%
  pivot_longer(c(home, away), names_to = 'location', values_to = 'team') %>%
  mutate(points = case_when(location == 'home' & sample_hgoal > sample_agoal ~ 3,
    location == 'away' & sample_agoal > sample_hgoal ~ 3,
    sample_hgoal == sample_agoal ~ 1,
    TRUE ~ 0)) %>%
  mutate(gd = case_when(location == 'home' ~ sample_hgoal - sample_agoal,
    location == 'away' ~ sample_agoal - sample_hgoal)))

#n = 10
#matchSims = rerun(n, nestedProb %>%
#  mutate(sampled_result = map(probabilities, #sample_n, 1, weight = prob)) %>%
#  select(-probabilities) %>%
#  unnest(cols = c(sampled_result)) %>%
#  select(-prob) %>%
#  pivot_longer(c(home, away), names_to = '#location', values_to = 'team') %>%
```

```
#           mutate(points = case_when(
#             location == 'home' & sample_hgoal > #sample_agoal ~ 3,
#             location == 'away' & sample_agoal > #sample_hgoal ~ 3,
#             sample_hgoal == sample_agoal ~ 1,
#             TRUE ~ 0
#           )) %>%
#           mutate(gd = case_when(
#             location == 'home' ~ sample_hgoal - #sample_agoal,
#             location == 'away' ~ sample_agoal - #sample_hgoal)))
```

From these simulations, I then take the average points earned for each team across a season, as well as their average goal differential to create a simulated Premier League table.

```
tableSim = matchSims %>%
  bind_rows() %>%
  group_by(team) %>%
  summarize(points = sum(points)/n,
             goalDiff = sum(gd)/n)

tableSim['team'][tableSim['team'] == 'Norwich'] = 'Nottingham'

finalTable = tableSim[order(-tableSim$points), ]
finalTable
```

```
## # A tibble: 20 x 3
##   team      points goalDiff
##   <chr>      <dbl>   <dbl>
## 1 Man City    91.6     70.4
## 2 Liverpool   84.4     53.7
## 3 Tottenham   69.8     27.7
## 4 Chelsea     69.4     25.9
## 5 Man United   67.0     22.8
## 6 Arsenal     63.0     16.6
## 7 Leicester   58.3      8.17
## 8 West Ham    51.5     -3.34
## 9 Wolves      50.3     -3.08
## 10 Aston Villa 48.8     -6.86
## 11 Brentford   48.1     -8.47
## 12 Everton     47.2     -9.35
## 13 Crystal Palace 45.7    -11.3
## 14 Leeds       43.8    -16.3
## 15 Newcastle   43.7    -14.4
## 16 Brighton    42.6    -14.8
## 17 Bournemouth 41.3    -20.1
## 18 Southampton 40.9    -20.5
## 19 Fulham      29.1    -38.1
## 20 Nottingham 19.0    -58.7
```

From this method, the top 4 teams are expected to be:

- Man City
- Liverpool

- Tottenham
- Chelsea